

Lecture Notes on Skip Lists


Lecture 12 — March 18, 2004

Erik Demaine

- Balanced tree structures we know at this point: B-trees, red-black trees, treaps.
- Could you implement them right now? Probably, with time... but without looking up any details in a book?
- Skip lists are a simple randomized structure you'll never forget.

Starting from scratch

- Initial goal: *just searches* — ignore updates (Insert/Delete) for now
- Simplest data structure: linked list
- Sorted linked list: $\Theta(n)$ time
- 2 sorted linked lists:
 - Each element can appear in 1 or both lists
 - How to speed up search?
 - **Idea:** Express and local subway lines
 - **Example:** 14, 23, 34, 42, 50, 59, 66, 72, 79, 86, 96, 103, 110, 116, 125
(What is this sequence?)
 - Boxed values are “express” stops; others are normal stops
 - Can quickly jump from express stop to next express stop, or from any stop to next normal stop
 - Represented as two linked lists, one for express stops and one for all stops:



- Every element is in linked list 2 (LL2); some elements also in linked list 1 (LL1)
- Link equal elements between the two levels
- To search, first search in LL1 until about to go too far, then go down and search in LL2

– Cost:

$$\text{len}(\text{LL1}) + \frac{\text{len}(\text{LL2})}{\text{len}(\text{LL1})} = \text{len}(\text{LL1}) + \frac{n}{\text{len}(\text{LL1})}$$

– Minimized when


$$\text{len}(\text{LL1}) = \frac{n}{\text{len}(\text{LL1})}$$

$$\Rightarrow \text{len}(\text{LL1})^2 = n$$

$$\Rightarrow \text{len}(\text{LL1}) = \sqrt{n}$$

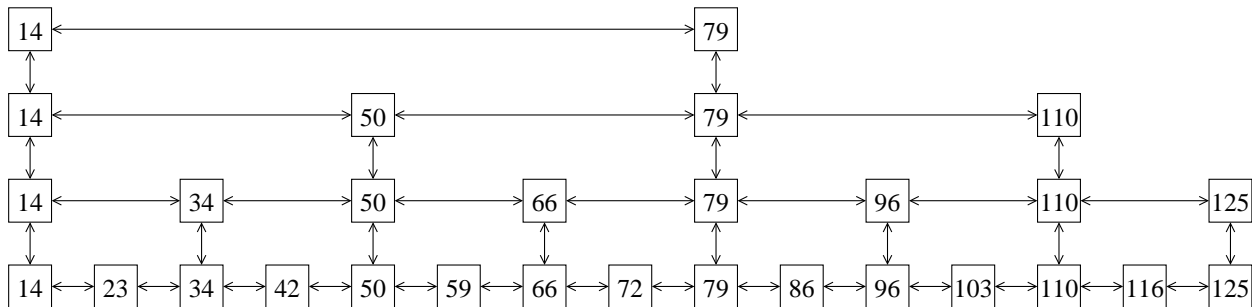
$$\Rightarrow \text{search cost} = 2\sqrt{n}$$

– Resulting 2-level structure:



- 3 linked lists: $3 \cdot \sqrt[3]{n}$
- k linked lists: $k \cdot \sqrt[k]{n}$
- $\lg n$ linked lists: $\lg n \cdot \sqrt[\lg n]{n} = \lg n \cdot \underbrace{n^{1/\lg n}}_{=2} = \Theta(\lg n)$

– Becomes like a binary tree:



– **Example:** Search for 72

- * Level 1: 14 too small, 79 too big; go down 14
- * Level 2: 14 too small, 50 too small, 79 too big; go down 50
- * Level 3: 50 too small, 66 too small, 79 too big; go down 66
- * Level 4: 66 too small, 72 spot on

Insert

- New element should certainly be added to bottommost level
(Invariant: Bottommost list contains all elements)
- Which other lists should it be added to?
(Is this the entire balance issue all over again?)
- **Idea:** Flip a coin
 - With what probability should it go to the next level?
 - To mimic a balanced binary tree, we'd like half of the elements to advance to the next-to-bottommost level
 - So, when you insert an element, flip a fair coin
 - If heads: add element to next level up, and flip another coin (repeat)
- Thus, on average:
 - $1/2$ the elements go up 1 level
 - $1/4$ the elements go up 2 levels
 - $1/8$ the elements go up 3 levels
 - Etc.
- Thus, “approximately even”

Example

- Get out a real coin and try an example
- You should put a special value $-\infty$ at the beginning of each list, and always promote this special value to the highest level of promotion
- This forces the leftmost element to be present in every list, which is necessary for searching

... many coins are flipped ...
(Isn't this easy?)

- The result is a skip list.
- It probably isn't as balanced as the ideal configurations drawn above.
- It's clearly good on average.
- Claim it's really really good, almost always.